

# RNA-seq analysis in R

## Contents

Disclaimer . . . . .	1
Setup . . . . .	1
Data . . . . .	1
Load data . . . . .	1
Formatting . . . . .	2
Filtering . . . . .	3
Annotations . . . . .	3
Visualization . . . . .	4
Correlations . . . . .	13
Differential expression . . . . .	16
GSEA . . . . .	22
Additional Material . . . . .	25
References . . . . .	30

## Disclaimer

The content of this tutorial is largely taken and adapted from the COMBINE RNA-seq Workshop, and from two tutorials using the DESeq2 package ( tutorial 1 and tutorial 2 ). We encourage you to visit these webpages to get extra material, including lecture slides and introductory R material if needed.

## Setup

If you haven't done so already, you should create a R project to store data, code, output tables and figures in separate sub-folders. Just click on "File > New Project..."

Create a new project. Add directories `./data`, `./R`, `./output` and `./figures`. Create a new R script file `script.R` in the `./R` directory. We will use it to copy paste useful command lines below. We can re-run all commands sequentially by clicking on the `source` button in the top right corner of the R studio editor.

## Data

The data for this tutorial comes from a Nature Cell Biology paper, *EGF-mediated induction of Mcl-1 at the switch to lactation is essential for alveolar cell survival* (Fu et al. 2015). Both the raw data (sequence reads) and processed data (counts) can be downloaded from Gene Expression Omnibus database (GEO) under accession number GSE60450.

This study examines the expression profiles of basal stem-cell enriched cells (B) and committed luminal cells (L) in the mammary gland of virgin, pregnant and lactating mice. Six groups are present, with one for each combination of cell type and mouse status. Each group contains two biological replicates. We will first use the counts file as a starting point for our analysis.

Data files are available from: <https://figshare.com/s/1d788fd384d33e913a2a>

Download the `SampleInfo_Corrected.txt` and `GSE60450_Lactation-GenewiseCounts.txt` files and place them in a `./data` directory within your project.

## Load data

We can now start by loading the data into the RStudio environment.

```

seqdata <- read.delim("./data/GSE60450_Lactation-GenewiseCounts.txt",
                      stringsAsFactors = FALSE)
sampleinfo <- read.delim("./data/SampleInfo_Corrected.txt")

```

You now have two objects in your environment. You can have a look at them using the `head()` or `View()` commands or by double clicking on an object within the environment. The `seqdata` object contains information about genes (one gene per row), the first column has the Entrez gene id, the second has the gene length and the remaining columns contain information about the number of reads aligning to the gene in each experimental sample.

```
names(seqdata)
```

```

## [1] "EntrezGeneID"
## [2] "Length"
## [3] "MCL1.DG_BC2CTUACXX_ACTTGA_L002_R1"
## [4] "MCL1.DH_BC2CTUACXX_CAGATC_L002_R1"
## [5] "MCL1.DI_BC2CTUACXX_ACAGTG_L002_R1"
## [6] "MCL1.DJ_BC2CTUACXX(CGATGT_L002_R1"
## [7] "MCL1.DK_BC2CTUACXX_TTAGGC_L002_R1"
## [8] "MCL1.DL_BC2CTUACXX_ATCACG_L002_R1"
## [9] "MCL1.LA_BC2CTUACXX_GATCAG_L001_R1"
## [10] "MCL1.LB_BC2CTUACXX_TGACCA_L001_R1"
## [11] "MCL1.LC_BC2CTUACXX_GCCAAT_L001_R1"
## [12] "MCL1.LD_BC2CTUACXX_GGCTAC_L001_R1"
## [13] "MCL1.LE_BC2CTUACXX_TAGCTT_L001_R1"
## [14] "MCL1.LF_BC2CTUACXX_CTTGTA_L001_R1"

```

The `sampleinfo` file contains basic information about the samples that we will need for the analysis today.

```
sampleinfo
```

	FileName	SampleName	CellType	Status
## 1	MCL1.DG_BC2CTUACXX_ACTTGA_L002_R1	MCL1.DG	basal	virgin
## 2	MCL1.DH_BC2CTUACXX_CAGATC_L002_R1	MCL1.DH	basal	virgin
## 3	MCL1.DI_BC2CTUACXX_ACAGTG_L002_R1	MCL1.DI	basal	pregnant
## 4	MCL1.DJ_BC2CTUACXX(CGATGT_L002_R1	MCL1.DJ	basal	pregnant
## 5	MCL1.DK_BC2CTUACXX_TTAGGC_L002_R1	MCL1.DK	basal	lactate
## 6	MCL1.DL_BC2CTUACXX_ATCACG_L002_R1	MCL1.DL	basal	lactate
## 7	MCL1.LA_BC2CTUACXX_GATCAG_L001_R1	MCL1.LA	luminal	virgin
## 8	MCL1.LB_BC2CTUACXX_TGACCA_L001_R1	MCL1.LB	luminal	virgin
## 9	MCL1.LC_BC2CTUACXX_GCCAAT_L001_R1	MCL1.LC	luminal	pregnant
## 10	MCL1.LD_BC2CTUACXX_GGCTAC_L001_R1	MCL1.LD	luminal	pregnant
## 11	MCL1.LE_BC2CTUACXX_TAGCTT_L001_R1	MCL1.LE	luminal	lactate
## 12	MCL1.LF_BC2CTUACXX_CTTGTA_L001_R1	MCL1.LF	luminal	lactate

## Formatting

We will be manipulating and reformatting the counts matrix into a suitable format for downstream analysis. The first two columns in the `seqdata` dataframe contain annotation information. We need to make a new matrix `countdata` containing only the counts, but we can store the gene identifiers (the `EntrezGeneID` column) as rownames.

```

# Remove first two columns from seqdata
countdata <- seqdata[,-(1:2)]
# Store EntrezGeneID as rownames
rownames(countdata) <- seqdata[,1]

```

The column names of `countdata` are the sample names which are pretty long so we'll shorten these to contain only the relevant information about each sample. We will use the `substr` command to extract the first 7 characters and use these as the colnames.

```
colnames(countdata) <- substr(colnames(countdata), start=1, stop=7)
```

Note that the column names are now the same as `SampleName` in the `sampleinfo` file. This is good because it means our sample information in `sampleinfo` is in the same order as the columns in `countdata`. Let's also simplify `sampleinfo` by putting `SampleName` as row names.

```
rownames(sampleinfo) <- sampleinfo[, 2]
sampleinfo <- sampleinfo[, -(1:2)]
table(colnames(countdata) == rownames(sampleinfo))
```

```
##  
## TRUE  
## 12
```

## Filtering

Genes with very low counts across all libraries provide little evidence for differential expression and they interfere with some of the statistical approximations that are used later in the pipeline. They also add to the multiple testing burden when estimating false discovery rates, reducing power to detect differentially expressed genes. These genes should be filtered out prior to further analysis. Here we perform a minimal pre-filtering to keep only genes that have at least 1 read total.

```
keep <- rowMeans(countdata) >= 1
countdata <- countdata[keep, ]
```

If you want to normalize data and perform differential expression analysis, you can jump to section Creating a `DESeqDataSet`

## Annotations

The only information we have about genes is their Entrez Gene ID, which is not very informative. We would like to add some annotation information. There are a number of ways to do this. We will demonstrate how to do this using the `org.Mm.eg.db` package.

First we need to decide what information we want. In order to see what we can extract we can run the `columns` function on the annotation database.

```
library(org.Mm.eg.db)
columns(org.Mm.eg.db)

## [1] "ACNUM"      "ALIAS"       "ENSEMBL"     "ENSEMLPROT"
## [5] "ENSEMLTRANS" "ENTREZID"    "ENZYME"      "EVIDENCE"
## [9] "EVIDENCEALL" "GENENAME"    "GO"          "GOALL"
## [13] "IPI"         "MGI"        "ONTOLOGY"    "ONTOLOGYALL"
## [17] "PATH"        "PFAM"       "PMID"        "PROSITE"
## [21] "REFSEQ"      "SYMBOL"     "UNIGENE"     "UNIPROT"
```

For now, let's get gene symbols and build up our annotation information in a separate data frame using the `select` function

```
annot <- select(org.Mm.eg.db, keys=rownames(countdata), columns=c("ENTREZID", "SYMBOL"))
```

Let's double check that the ENTREZID column matches exactly to our countdata rownames.

```
table(annot$ENTREZID==rownames(countdata))
```

```
##  
##   TRUE  
## 18178
```

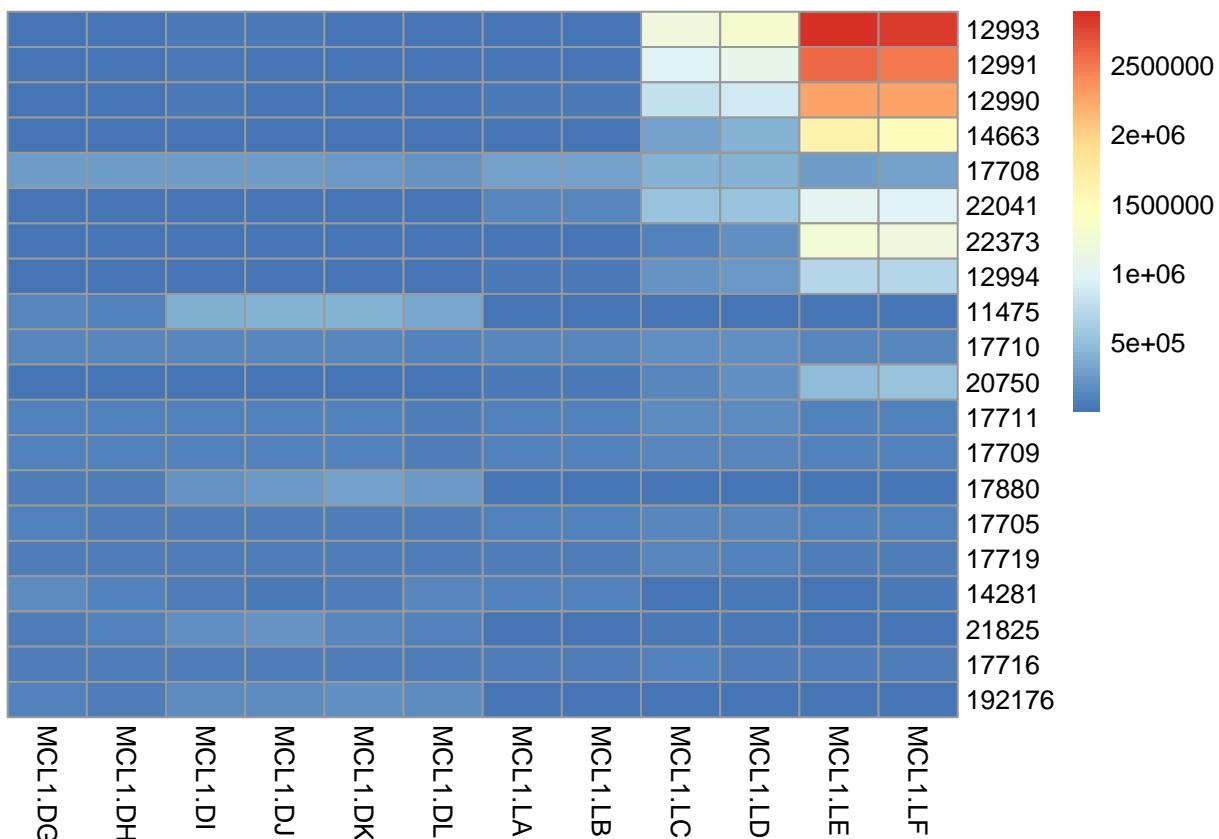
## Visualization

We are working we raw count data in this section. Data normalization and transformation will come later.

### Heatmap

To explore a count matrix, it is often instructive to look at it as a heatmap. Below we show how to produce such a heatmap. We focus on the 20 genes with the highest average counts.

```
library("pheatmap")  
  
select <- order(rowMeans(countdata), decreasing=TRUE) [1:20]  
  
pheatmap(countdata[select, ],  
         cluster_rows=FALSE,  
         cluster_cols=FALSE)
```



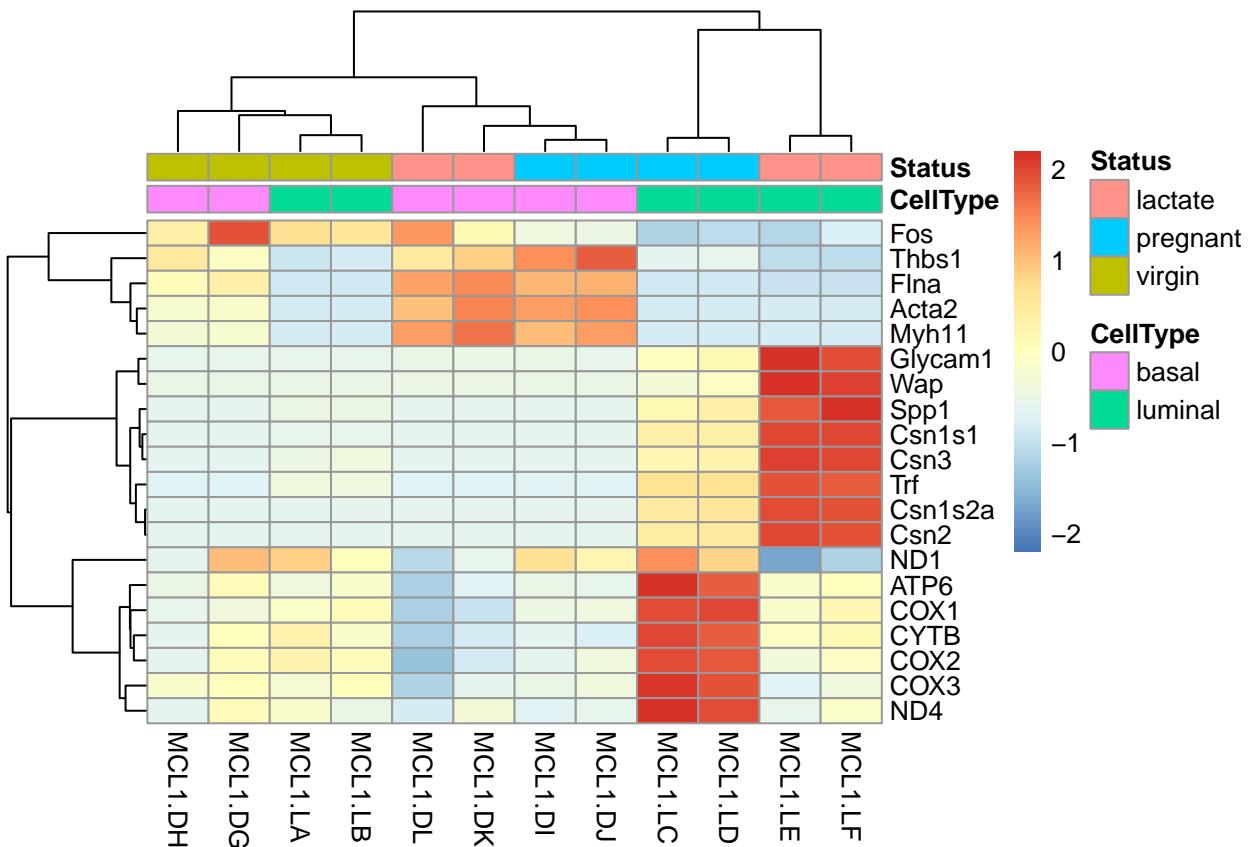
Look up the documentation for `pheatmap` in Rstudio. The parameter `annotation_col` allow us to add annotation columns. We'll also cluster rows and columns , scale values by rows and change row labels using gene symbol.

```
pheatmap(countdata[select, ],  
         show_rownames=TRUE,
```

```

    cluster_cols=TRUE,
    annotation_col=sampleinfo[ , c("CellType","Status")] ,
    scale = "row",
    labels_row = annot$SYMBOL[select]
)

```



If you like interactive objects, you could try the `heatmaply()` function to generate the heatmap.

```

library(heatmaply)

heatmaply(countdata[select,], scale = "row")

```

## PCA

Principal component analysis (PCA) is a great tool to see the overall “shape” of the data. It allows to identify which samples are similar to one another and which are very different. This can enable us to identify groups of samples that are similar and work out which variables make one group different from another. We use the `prcomp` function to run the PCA. Usually, features (here genes) are columns while observation (here samples) are rows so we'll transpose `countdata` to before running the PCA.

```

pca_res <- prcomp( t(countdata), center = TRUE, scale. = TRUE)
summary(pca_res)

```

```

## Importance of components:
##                               PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation     83.8370 72.0165 40.90397 38.68490 31.0671 22.45107
## Proportion of Variance 0.3867 0.2853 0.09204 0.08233 0.0531 0.02773
## Cumulative Proportion   0.3867 0.6720 0.76401 0.84633 0.8994 0.92716

```

```

##          PC7      PC8      PC9      PC10     PC11
## Standard deviation   18.77350 17.49166 16.17556 15.66165 12.60086
## Proportion of Variance 0.01939  0.01683  0.01439  0.01349  0.00873
## Cumulative Proportion 0.94655  0.96338  0.97777  0.99127  1.00000
##          PC12
## Standard deviation   1.652e-13
## Proportion of Variance 0.000e+00
## Cumulative Proportion 1.000e+00

```

You obtain 12 principal components, each one explaining a percentage of the total variation in the dataset (PC1 explains 39%, PC2 29% and so on). The relationship (correlation or anticorrelation, a.k.a loadings) between the initial variables and the principal components is in `$rotation`. Let's see which genes have the greatest loadings onto PC1:

```

idx_up <- order(pca_res$rotation[, 1], decreasing = TRUE)
idx_down <- order(pca_res$rotation[, 1], decreasing = FALSE)
head(pca_res$rotation[idx_up, 1:2])

```

```

##          PC1      PC2
## 214531 0.01182908 0.0006886089
## 22778  0.01176338 0.0001653767
## 114304 0.01174112 0.0010381224
## 142980 0.01173840 -0.0003178144
## 66536  0.01171238 -0.0003867467
## 69727  0.01170085  0.0012227270

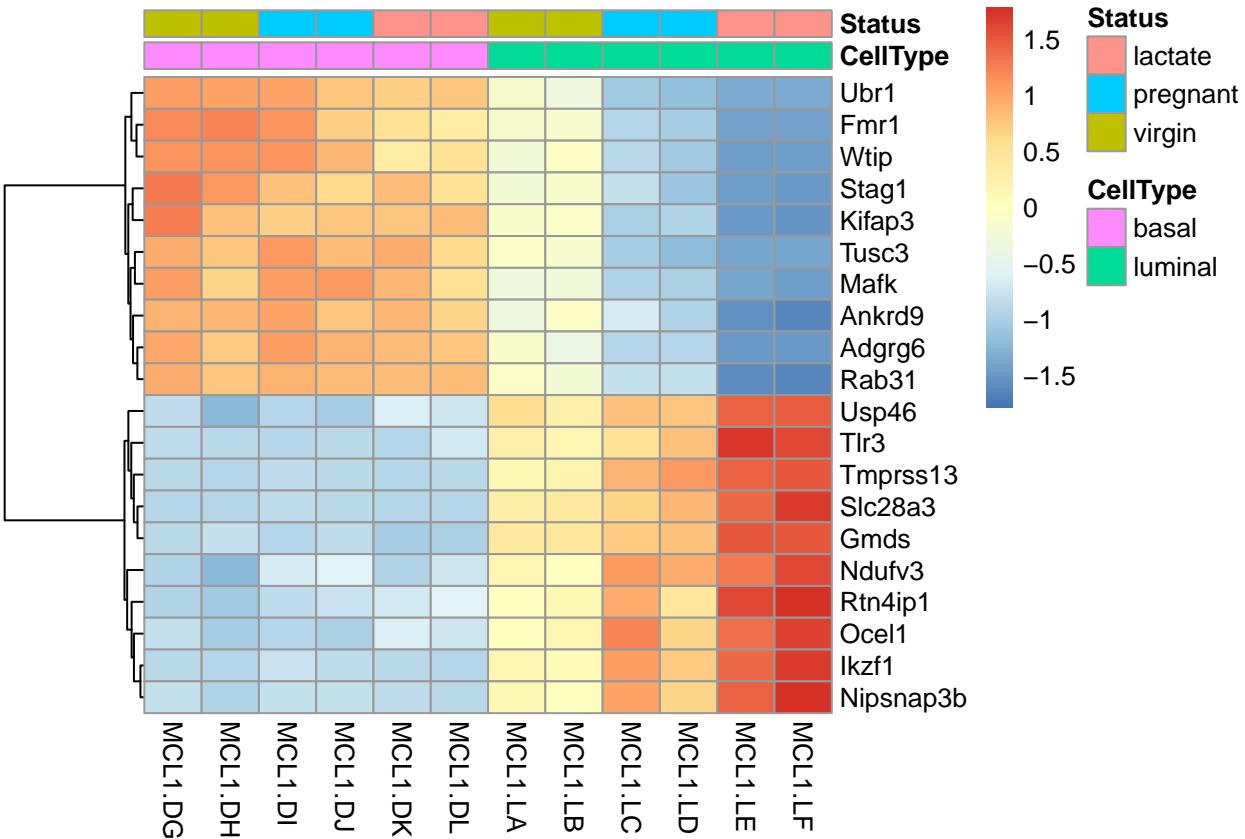
```

Let's look at the count data heatmap for genes with greatest and lowest loadings onto PC1. How does it look for PC2?

```

idx_gene <- c(idx_up[1:10], idx_down[1:10])
pheatmap(countdata[idx_gene, ],
         show_rownames=TRUE,
         cluster_cols=FALSE,
         annotation_col=sampleinfo[, c("CellType", "Status")],
         scale = "row",
         labels_row = annot$SYMBOL[idx_gene]
)

```



The values of each sample in terms of the principal components is in `$x`. Check taht the row names `pca_res$x` are the same as that of `sampleinfo`:

```
pca_res$x[, 1:2]
```

```
##          PC1        PC2
## MCL1.DG -91.91003  21.174002
## MCL1.DH -79.97473 -2.987728
## MCL1.DI -82.15576 -12.811208
## MCL1.DJ -67.73988 -40.945893
## MCL1.DK -61.12426 -67.292207
## MCL1.DL -55.00054 -67.320751
## MCL1.LA  17.84360 117.724564
## MCL1.LB  12.05523 137.439876
## MCL1.LC  75.36071  44.429079
## MCL1.LD  80.94125  16.790115
## MCL1.LE 125.65812 -73.007205
## MCL1.LF 126.04628 -73.192642
```

We'll use the `ggplot2` package to plot the results of the PCA. See the R for data science book for an introduction to `ggplot2`. We first need to create a data frame from the pca results.

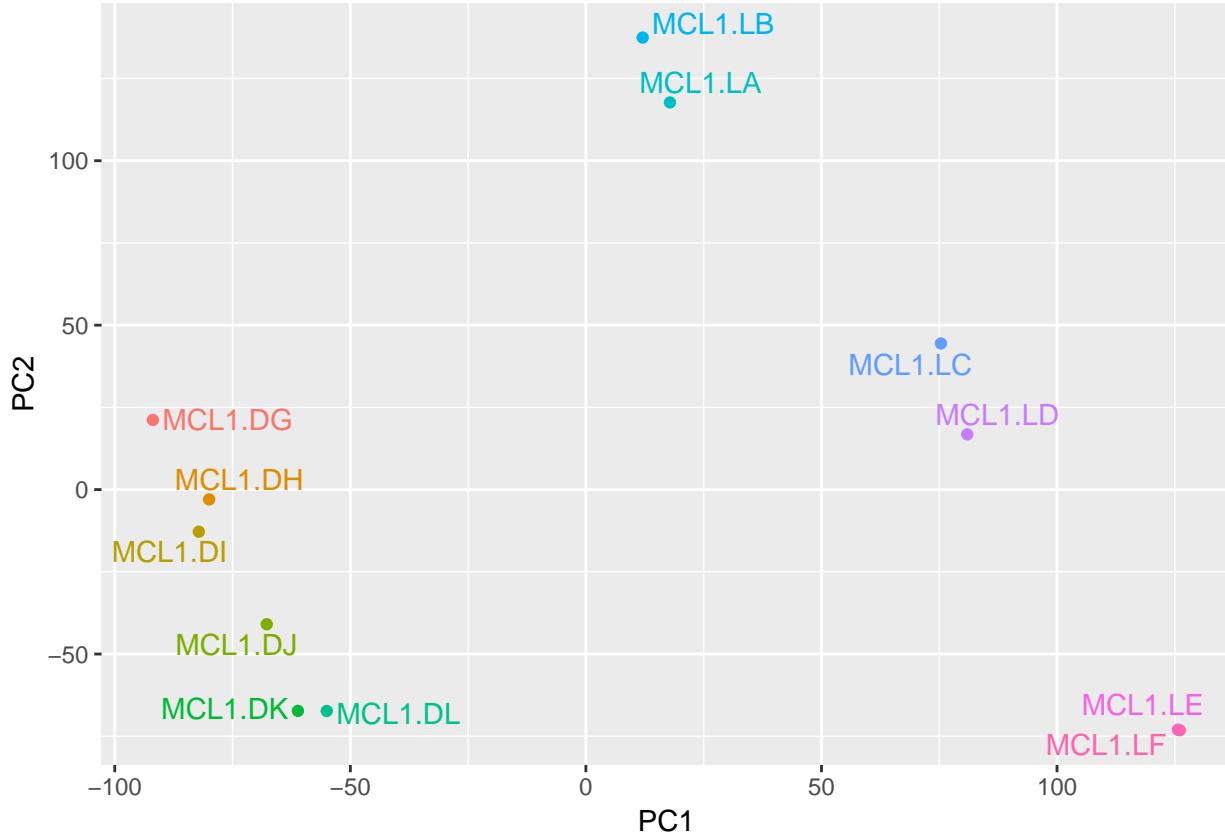
```
df_pca <- as.data.frame(pca_res$x)
df_pca$sample <- rownames(df_pca)
names(df_pca)

## [1] "PC1"      "PC2"      "PC3"      "PC4"      "PC5"      "PC6"      "PC7"
## [8] "PC8"      "PC9"      "PC10"     "PC11"     "PC12"     "sample"
```

We can now use the `ggplot()` function and choose to draw one point per sample.

```
library(ggplot2)
library(ggrepel)

pca_plot <- ggplot(df_pca, aes(x=PC1, y=PC2, color = sample, label=sample)) +
  geom_point(show.legend = FALSE) +
  geom_text_repel(show.legend = FALSE)
pca_plot
```



As an exercise, check that the row names `pca_res$x` are the same as that of `sampleinfo` (check it) and map metadata directly to the `df_pca` dataframe. Color points according to the sample Status or CellType.

->

Do the same using t-SNE

```
library(Rtsne)
tsne_res <- Rtsne(t(countdata), perplexity = 3)
df_tsne <- as.data.frame(tsne_res$Y)
colnames(df_tsne) <- c("tSNE1", "tSNE2")
df_tsne$sample <- colnames(countdata)

ggplot(df_tsne, aes(x=tSNE1, y=tSNE2, color = sample, label=sample)) +
  geom_point(alpha = 0.25, show.legend = FALSE) +
  geom_text_repel(show.legend = FALSE)
```

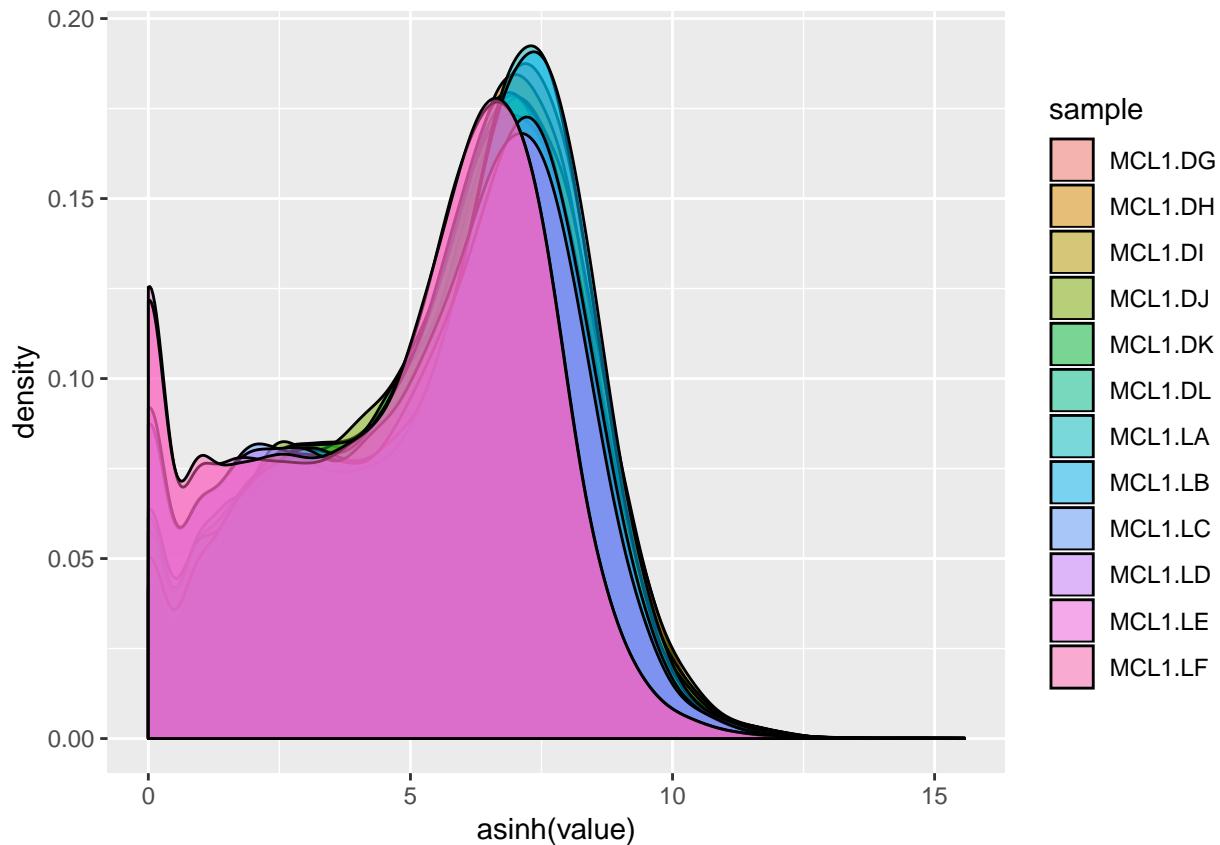
## Manipulating the data

Data visualization with `ggplot2` works with data frames. Here, we convert our data into a data frame where all count values are stored in the same column named `value` (long format). This will allow us to fully exploit `ggplot2` features.

```
library(reshape2)
library(tidyverse)
df <- countdata
df$GeneID <- rownames(countdata)
df_melt <- melt(df, id.vars = "GeneID")
df_melt <- rename(df_melt, sample = variable)
```

As an example, we can now plot the distribution of asinh transformed count values for each sample (You can offset histograms for better readability using the `ggridges` package).

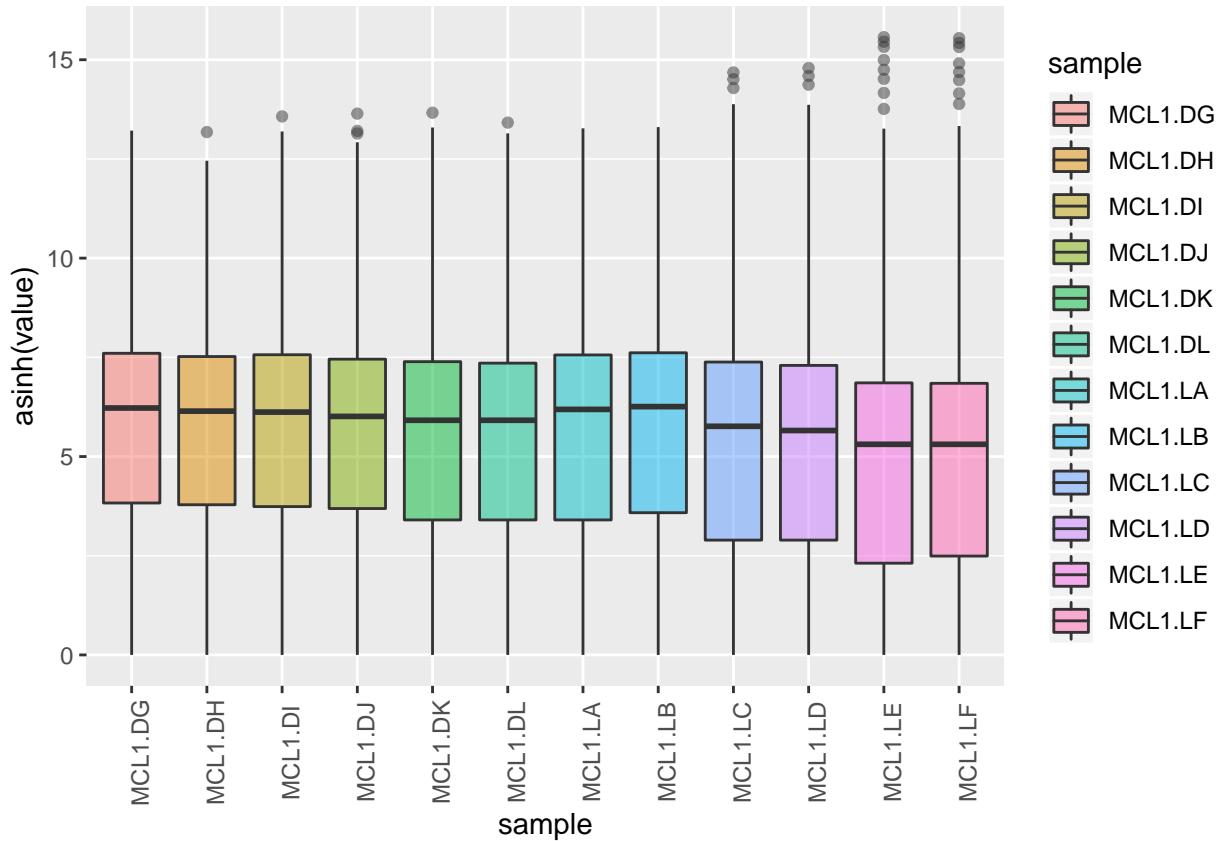
```
plot <- ggplot(df_melt, aes(x=asinh(value), fill = sample)) +
  geom_density(alpha = 0.5)
plot
```



Box plots allow to see a summary of these distributions.

```
library(ggridges)

plot <- ggplot(df_melt, aes(x=sample, y = asinh(value), fill = sample)) +
  theme(axis.text.x = element_text(angle = 90)) +
  geom_boxplot(alpha = 0.5)
plot
```



The last two samples seem to differ quite neatly from the other samples. Some normalization between samples will be needed before going further into the analysis.

Here we use the `dplyr` package to group rows and return group summary (see the corresponding chapter in R for data science ). It makes it easy to compute the median count value per sample.

```
df_median <-
  df_melt %>%
  group_by(sample) %>%
  summarise(median = median(value, na.rm = TRUE))

df_median

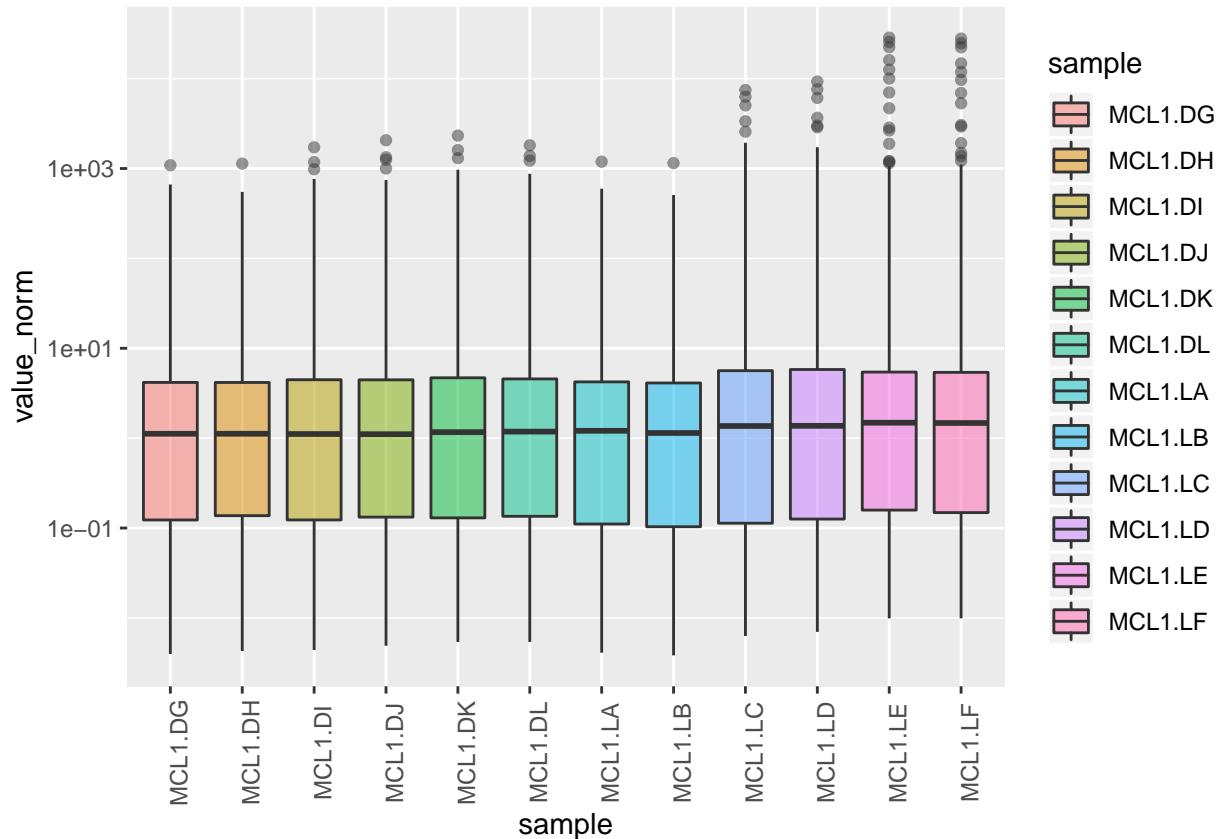
## # A tibble: 12 x 2
##   sample  median
##   <fct>    <dbl>
## 1 MCL1.DG  252
## 2 MCL1.DH  233
## 3 MCL1.DI  228.
## 4 MCL1.DJ  204
## 5 MCL1.DK  185
## 6 MCL1.DL  185
## 7 MCL1.LA  244.
## 8 MCL1.LB  261
## 9 MCL1.LC  159
## 10 MCL1.LD 143
## 11 MCL1.LE 101
## 12 MCL1.LF 101
```

Build a data frame with the median and mean count per gene across samples.

Normalization will be carried out properly later. Here we show how we can use the `dplyr` package to normalize the data using the median:

```
df_melt_norm <-
  df_melt %>%
  group_by(sample) %>%
  mutate(value_norm = value/median(value, na.rm = TRUE))

ggplot(df_melt_norm, aes(x=sample, y = value_norm, fill = sample)) +
  theme(axis.text.x = element_text(angle = 90)) +
  geom_boxplot(alpha = 0.5) +
  scale_y_log10()
```



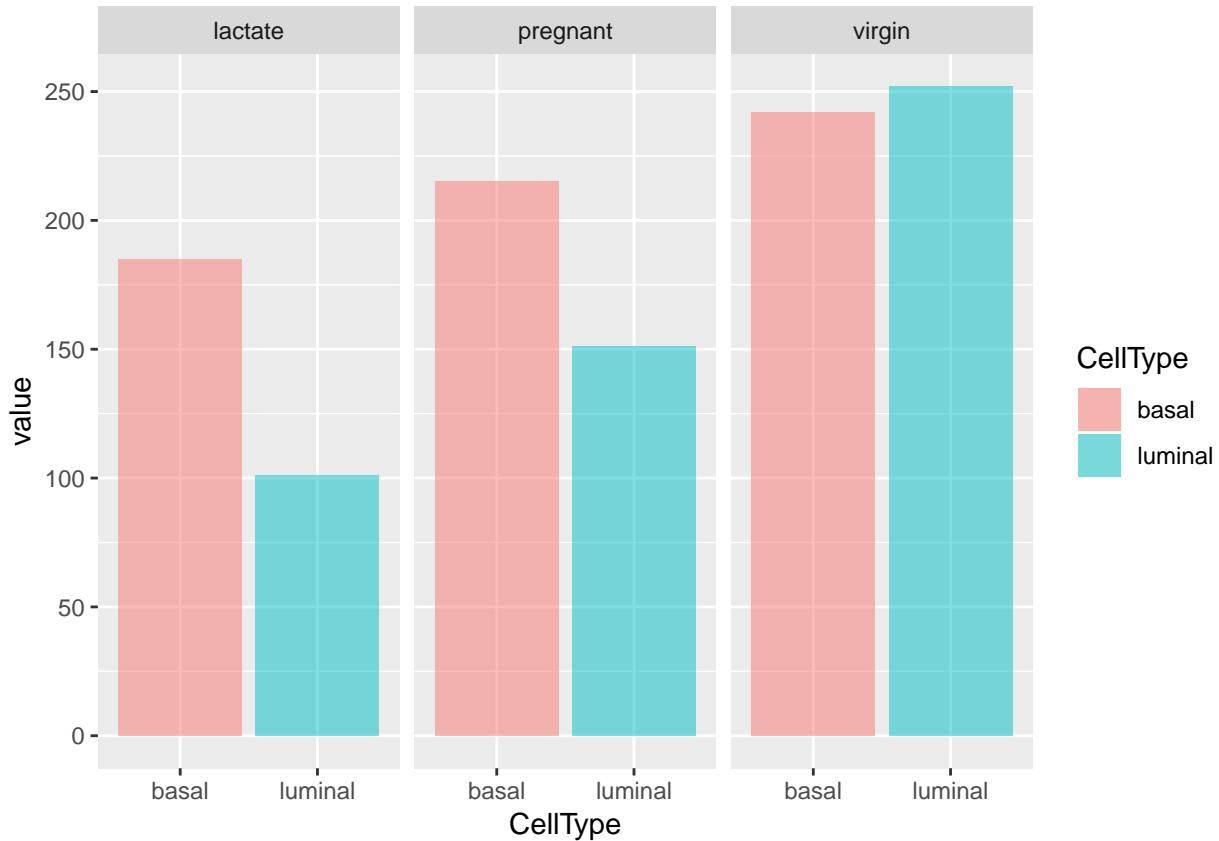
## Using metadata

Things get more interesting when we include metadata. Let's map sample information to `df_melt`

```
idx_match <- match(df_melt$sample, rownames(sampleinfo))
df_melt <- cbind(df_melt, sampleinfo[idx_match, ])
```

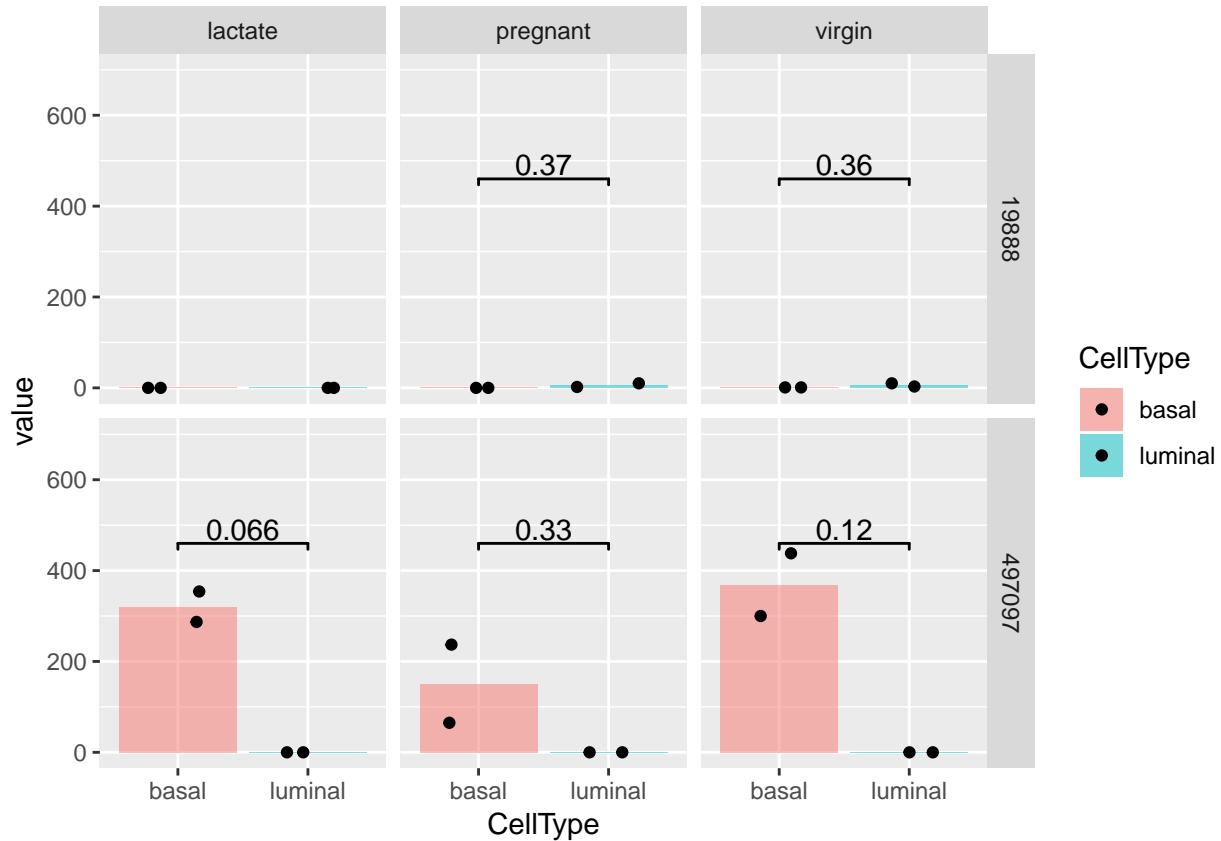
Now we have access to two more variables, `CellType` and `Status`, that we could use to form groups.

```
ggplot(df_melt, aes(x = CellType, y = value, fill = CellType)) +
  geom_bar(alpha = 0.5, stat = "summary", fun.y = "median", position = "dodge") +
  facet_wrap(~Status)
```



We can also focus on a particular set of GeneIDs and split plots using `facet_grid`. Let's add an unpaired t-test on top of that.

```
library(ggsignif)
idx_select <- df_melt$GeneID %in% df_melt$GeneID[1:2]
ggplot(df_melt[idx_select, ], aes(x = CellType, y = value, fill = CellType)) +
  geom_bar(alpha = 0.5, stat = "summary", fun.y = "median", position = "dodge") +
  geom_point(position = position_jitter(width = 0.25, height = 0))+
  geom_signif(comparisons = list(1:2),
              na.rm = TRUE, test = "t.test",
              test.args = list("paired" = FALSE),
              position = "identity") +
  ylim(c(0,700))+
  facet_grid(GeneID~Status)
```



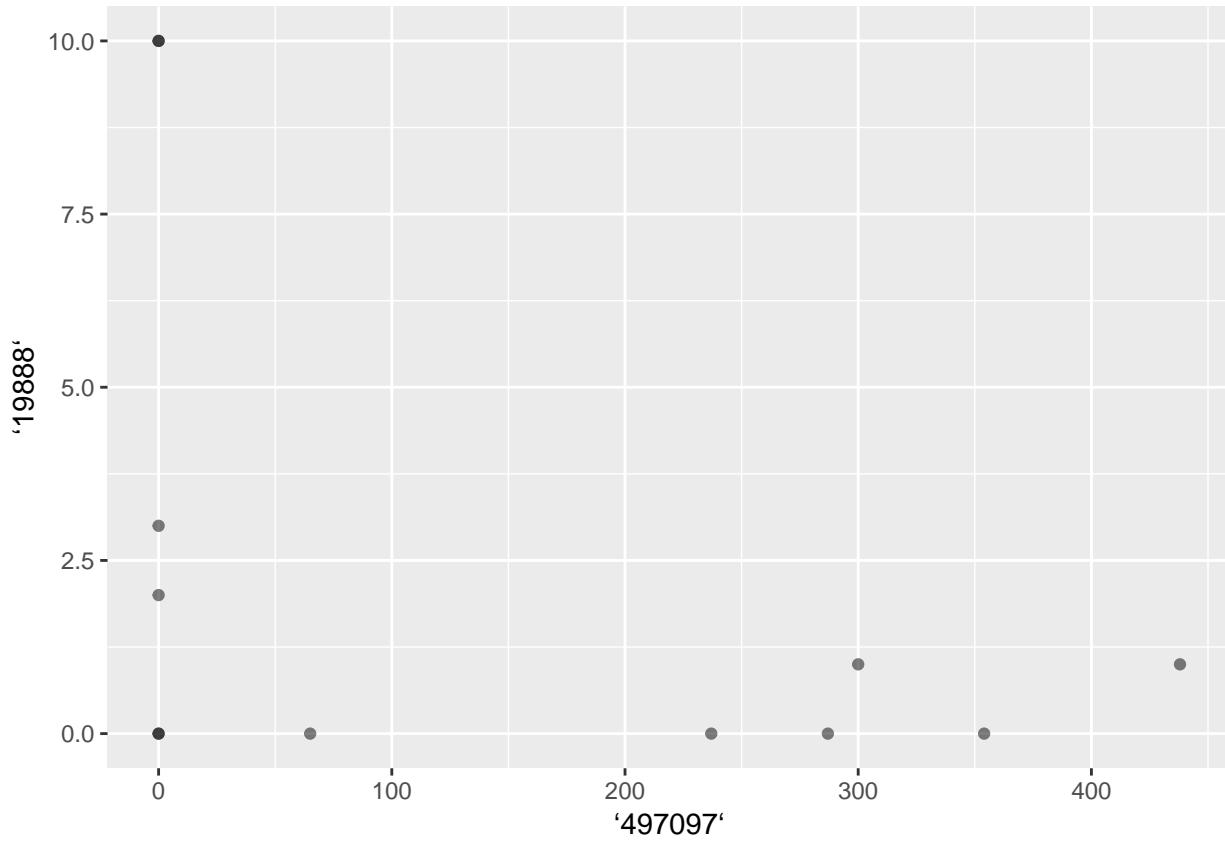
## Correlations

At some point in the analysis, we might want to perform pairwise comparisons between genes. Let's first create a data frame (note that we use the transpose of the countdata matrix so available variables are in `names(df)`).

```
df <- as.data.frame(t(countdata))
```

Plotting the data for two different genes is a good way to identify functional relationships between them. Let's pick two genes,

```
library(ggplot2)
xvar <- as.name(names(df)[1])
yvar <- as.name(names(df)[2])
p <- ggplot(df, aes_string(x=xvar, y=yvar)) + geom_point(alpha = 0.5)
p
```



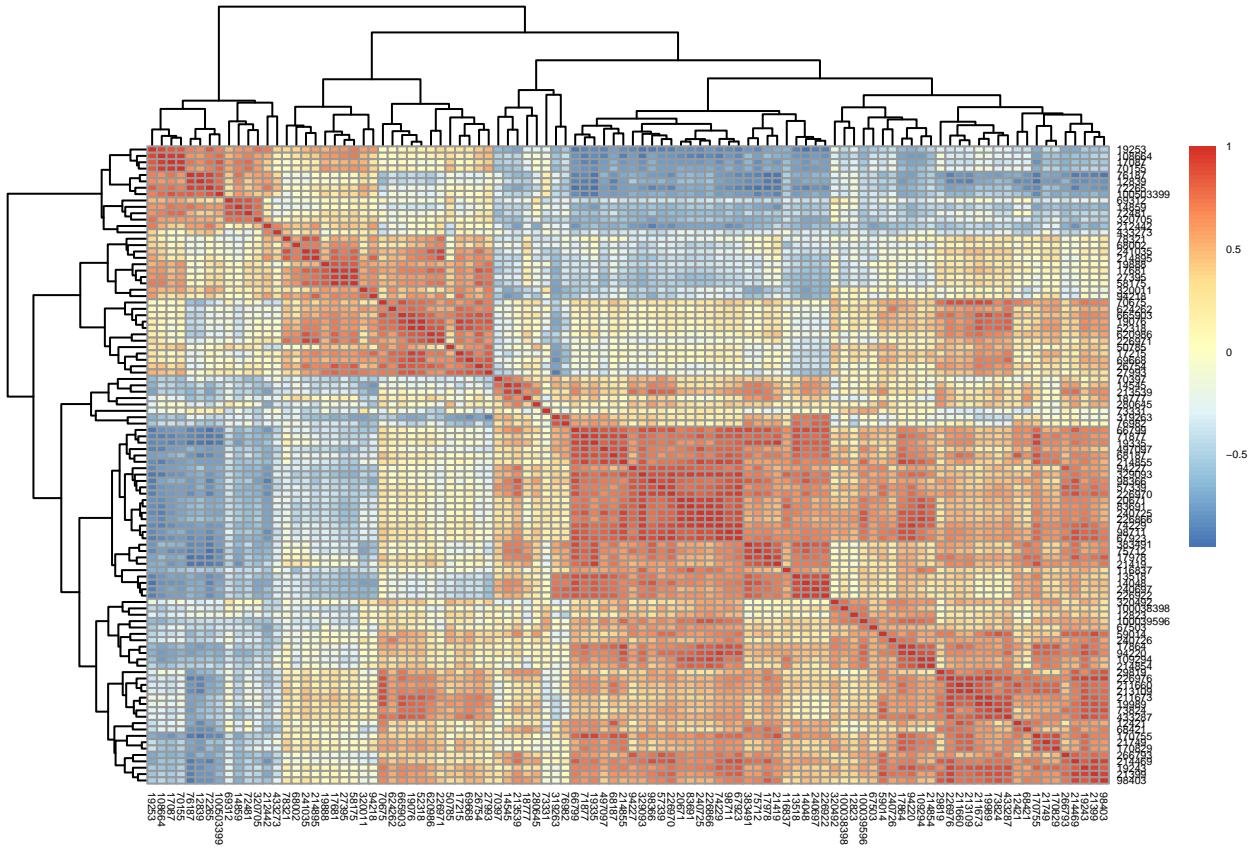
and fit a linear model to the data and compute Pearson's correlation coefficient.

```
lm_res <- lm(formula = `19888` ~ `497097`, data = df)
sm<-summary.lm(lm_res)
sqrt(sm$r.squared)
```

```
## [1] 0.4371161
```

Now we want to repeat this for all pairs of genes (we use only the 100 first genes here) and build a correlation matrix

```
library(Hmisc)
corr <- Hmisc::rcorr(t(countdata[1:100, ]))
pheatmap(corr$r, fontsize = 4)
```



We can explore correlations by constructing a data frame with the correlation coefficient, its associated p-value and the number of points used to compute the correlation

```
library(reshape2)
df_R_coeff <- melt(corr$r)
df_R_coeff$variable <- "R"
df_pval <- melt(corr$p)
df_pval$variable <- "P"
df_nval <- melt(corr$n)
df_nval$variable <- "n"

df_corr <- rbind(df_R_coeff, df_pval, df_nval)
df_corr <- dcast(df_corr, Var1 + Var2 ~ variable, mean)
summary(df_corr)
```

	Var1	Var2	n	P
##	Min. : 12421	Min. : 12421	Min. : 12	Min. : 0.00000
##	1st Qu.: 27844	1st Qu.: 27844	1st Qu.: 12	1st Qu.: 0.01726
##	Median : 75950	Median : 75950	Median : 12	Median : 0.13106
##	Mean : 3140498	Mean : 3140498	Mean : 12	Mean : 0.25801
##	3rd Qu.: 226880	3rd Qu.: 226880	3rd Qu.: 12	3rd Qu.: 0.43326
##	Max. : 100503399	Max. : 100503399	Max. : 12	Max. : 0.99946
##			NA's : 100	
##	R			
##	Min. :-0.9476			
##	1st Qu.:-0.2527			
##	Median : 0.2527			
##	Mean : 0.1753			

```
## 3rd Qu.: 0.6052
## Max. : 1.0000
##
```

Now we can filter and keep only the most correlated pairs

```
df_corr_filtered <- df_corr[!is.na(df_corr$P) & df_corr$n>10 & df_corr$P<0.05 & df_corr$R>0.9, ]
```

We use that restricted set to build a network.

```
library(igraph)

net <- igraph::graph.data.frame(df_corr_filtered[, c("Var1", "Var2")], directed=FALSE)
net <- igraph::simplify(net)
clusters <- igraph::cluster_fast_greedy(as.undirected(net))
group = clusters$membership
names(group) <- clusters$names
```

And we render it in an interactive environment

```
library(networkD3)
net_d3 <- networkD3::igraph_to_networkD3(net, group = group)
p <- forceNetwork(Links = net_d3$links, Nodes = net_d3$nodes,
                  Source = 'source', Target = 'target',
                  fontFamily = "arial",
                  NodeID = 'name', Group = 'group',
                  colourScale = JS("d3.scaleOrdinal(d3.schemeCategory20);"),
                  charge = -10, opacity = 1,
                  linkColour = rgb(0.75, 0.75, 0.75),
                  fontSize = 12, bounded = TRUE, zoom=TRUE, opacityNoHover = 1)
p
```

Alternatively, we can save it as a text file and use another software such as *cytoscape* to manipulate this network.

```
dir.create("./output/")
write.table(df_corr_filtered, file = "./output/df_corr_filtered.txt", sep = "\t", row.names = FALSE, quote = FALSE)
write.table(data.frame(clusters$membership, clusters$names), file = "./output/group.txt", sep = "\t", quote = FALSE)
```

Exercise: add an extra column df\$score <- 1:length(rownames(df)) that contains a score associated to each sample to our data frame. Which genes most strongly correlate with this score?

## Differential expression

### Creating a DESeqDataSet

We will use the DESeq2 package to conduct the core steps of the analysis. It should be already installed so we just need to load it in the RStudio session.

```
library(DESeq2)
```

We can use the `help("DESeq2-package")` command to browse the package documentation in Rstudio. We will use the `DESeqDataSetFromMatrix` function to build a `DESeqDataSet` object.

```
dds <- DESeqDataSetFromMatrix(countData = countdata,
                               colData = sampleinfo,
                               design = ~ CellType + Status)
```

The design indicates how to model the samples, i.e how the counts for each gene depend on the variables in

`colData = sampleinfo`. Here we want to measure the effect of the cell type and of the pregnancy status of the mice. Note that the two factor variables `CellType` and `Status` should be columns of `sampleinfo`. Now the `dds` object contains count data along with the metadata and the experiment design. The count data is obtained using `counts(dds)` and the metadata is obtained using `colData(dds)`.

Now that we have a `DESeqDataSet` object, we can analyse the data using the many tools available in the `DESeq2` package.

### Differential expression analysis

The standard differential expression analysis steps are wrapped into a single function, `DESeq`. It is then straightforward to perform differential expression analysis.

```
dds <- DESeq(dds)
```

We can see the comparisons carried out by `DESeq` using `resultsNames()`.

```
resultsNames(dds)
```

```
## [1] "Intercept"                  "CellType_luminal_vs_basal"
## [3] "Status_pregnant_vs_lactate" "Status_virgin_vs_lactate"
```

Results tables are generated using the function `results`, which extracts a results table with log2 fold changes, p values and adjusted p values. We specify which comparison we want to extract using the `name` parameter. Here we choose to compare basal and luminal cell types. We select significant results with adjusted p-values (correction for multiple tests computed with the Benjamini–Hochberg procedure) below 0.01 using `alpha`.

```
res <- results(dds, alpha = 0.001, pAdjustMethod="BH", name="CellType_luminal_vs_basal")
head(res)
```

```
## log2 fold change (MLE): CellType luminal vs basal
## Wald test p-value: CellType luminal vs basal
## DataFrame with 6 rows and 6 columns
##           baseMean      log2FoldChange       lfcSE
##           <numeric>      <numeric>       <numeric>
## 497097 126.481585433975 -10.1820437902145 0.930694407261515
## 19888  1.95697709143844   3.067152056161  1.23908077258863
## 20671  50.5979096011351  -2.71203870035955 0.337521822867245
## 27395  358.719720532762   0.799404687904274  0.127971895175
## 18777  744.356910502922  0.106722300367076  0.134155732874136
## 21399  1097.91010559132 -0.361177585114161  0.0647349104833116
##           stat        pvalue       padj
##           <numeric>      <numeric>       <numeric>
## 497097 -10.9402653661305 7.39830425247185e-28 9.63986354984643e-27
## 19888   2.47534472652113  0.0133107627365738  0.0207548596551985
## 20671  -8.03515066765403  9.34639809275225e-16 5.95803042227874e-15
## 27395   6.24672071013012  4.19159576750842e-10 1.67328656261537e-09
## 18777   0.795510546442337  0.42631657857844  0.482508269652693
## 21399  -5.57933242538849  2.41443422519718e-08 8.05872191344228e-08
```

We can get a summary of these results using `summary()`

```
summary(res)
```

```
##
## out of 18178 with nonzero total read count
## adjusted p-value < 0.001
## LFC > 0 (up)      : 4185, 23%
## LFC < 0 (down)    : 4741, 26%
```

```

## outliers [1]      : 0, 0%
## low counts [2]    : 705, 3.9%
## (mean count < 2)
## [1] see 'cooksCutoff' argument of ?results
## [2] see 'independentFiltering' argument of ?results

Let's select genes have an adjusted p-value lower than 0.001 and a fold-change greater than 2^5:

enrich <- data.frame(ID=rownames(res), res@listData)

enrich <- enrich %>%
  filter(padj < 0.001, log2FoldChange > 5) %>%
  arrange(desc(log2FoldChange))

enrich %>% head(10)

##           ID baseMean log2FoldChange      lfcSE      stat     pvalue
## 1  192200  510.96091      7.836400 0.6605401 11.863625 1.828812e-32
## 2  791407   72.08411      7.785638 0.7407422 10.510590 7.720906e-26
## 3  21990   977.00600      7.686764 0.5023793 15.300719 7.561686e-53
## 4  217316   58.82642      7.482389 0.7202352 10.388813 2.788000e-25
## 5  382036   16.46267      7.472099 0.8865860  8.427946 3.517839e-17
## 6  12310   450.32143      7.388458 0.6800650 10.864340 1.704520e-27
## 7  20309   366.10363      7.296678 0.4709648 15.493044 3.865462e-54
## 8  246133   16.26991      7.047306 1.1689753  6.028618 1.653672e-09
## 9   71601   27.35394      7.010957 0.8244734  8.503558 1.838672e-17
## 10 100503562  99.14758      6.949689 0.4788576 14.513059 1.001526e-47
##          padj
## 1  3.011766e-31
## 2  9.183620e-25
## 3  2.735514e-51
## 4  3.202809e-24
## 5  2.457705e-16
## 6  2.185112e-26
## 7  1.471486e-52
## 8  6.262377e-09
## 9  1.312919e-16
## 10 3.108289e-46

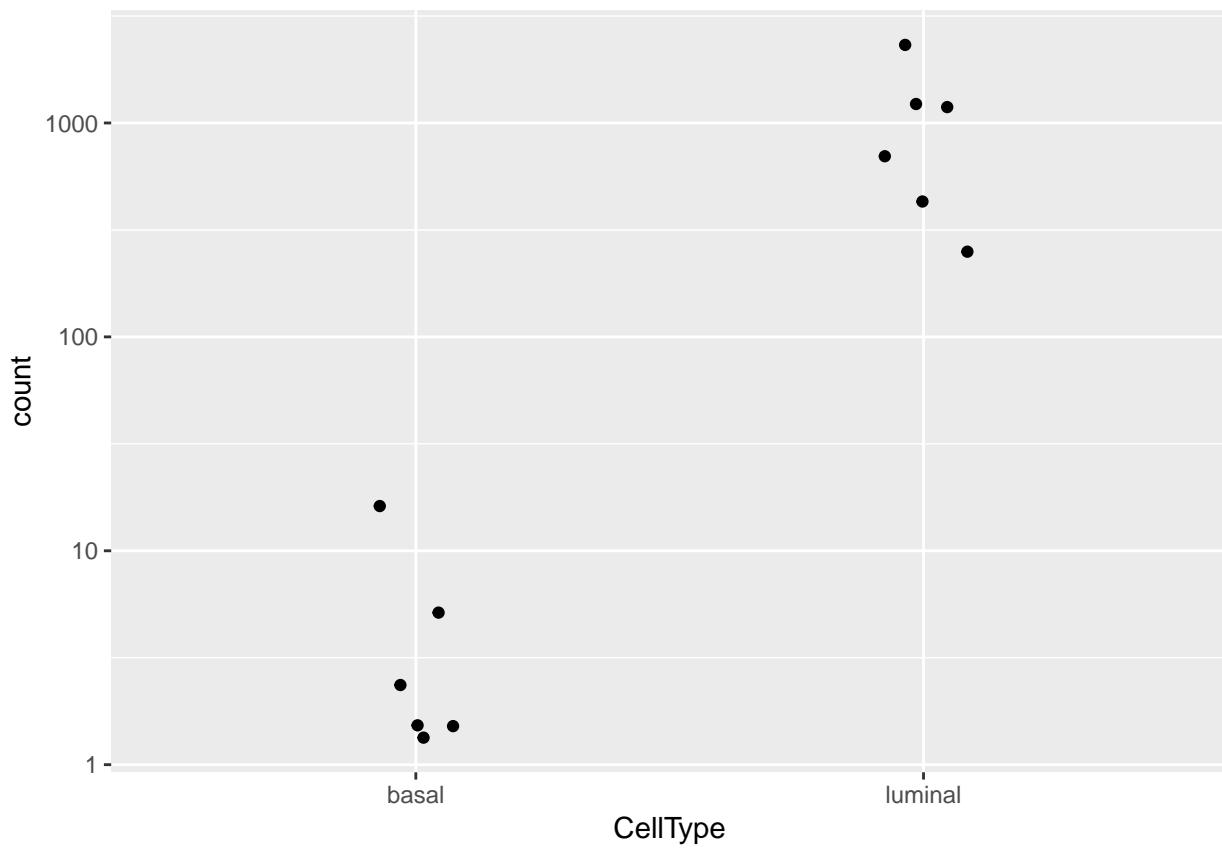
```

Have a look at the most enriched gene

```

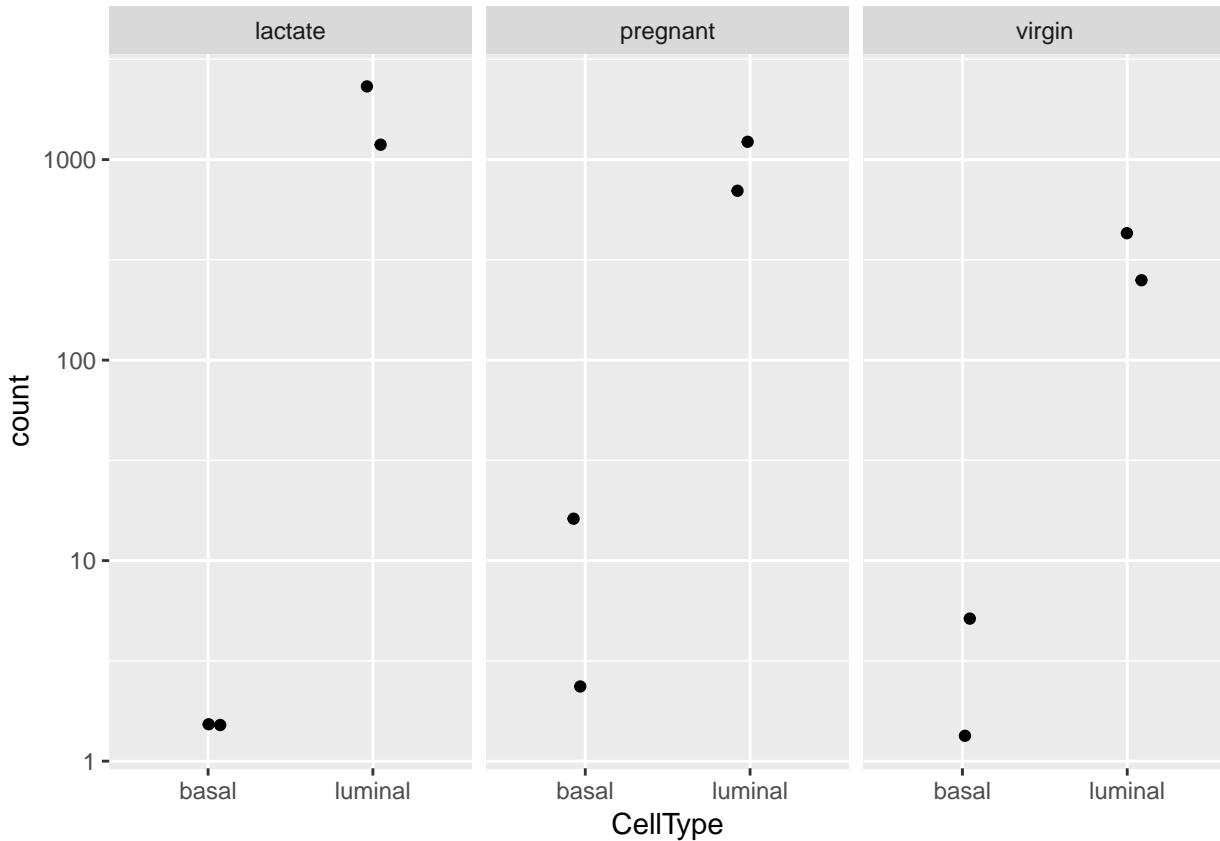
d <- plotCounts(dds, gene=as.character(enrich$ID[1]), intgroup=c("CellType","Status"),
                 returnData=TRUE)
library("ggplot2")
p <- ggplot(d, aes(x=CellType, y=count)) +
  geom_point(position=position_jitter(w=0.1,h=0)) +
  scale_y_log10()
p

```



We can also facet the plot according to `Status`. It could be interesting to see how the `CellType` effect depends on `Status`. See section Interactions to address such questions.

```
p + facet_wrap(~Status)
```



We have used default parameters here. Data normalization and model fitting have been performed in the background by DESeq. If you want to perform these steps separately, have a look at the Additional Material section.

## Interactions

Interaction terms allow to test, for example, if the log<sub>2</sub> fold change attributable to a given condition is different based on another factor, for example if the `CellType` effect differs across `Status`.

Interaction terms can be added to the design formula but a simpler approach to take into account interactions consists in:

- combining the factors of interest into a single factor with all combinations of the original factors
- changing the design to include just this factor, e.g. `~ group`

```
dds.int <- dds
dds.int$group <- factor(paste(dds.int$status, dds.int$CellType, sep="."))
design(dds.int) <- ~ group
dds.int <- DESeq(dds.int)
```

Now we can compare basal and luminal cell types within the lactate status.

```
res.int <- results(dds.int, contrast = c("group", "lactate.basal", "lactate.luminal"))
res.int

## log2 fold change (MLE): group lactate.basal vs lactate.luminal
## Wald test p-value: group lactate.basal vs lactate.luminal
## DataFrame with 18178 rows and 6 columns
##           baseMean     log2FoldChange      lfcSE
##           <numeric>       <numeric>       <numeric>
## 497097    126.481585433975   9.99193033684966  1.70778319837909
## 19888     1.95697709143844          0  2.74448002297715
```

```

## 20671      50.5979096011351   2.49566717862655  0.627259620835026
## 27395      358.719720532762 -0.986661733396428  0.201363861143128
## 18777      744.356910502922 -0.1627472423715  0.141156023491652
## ...
## ...
## 100862004  3409.10829251673 -0.114306553961534  0.19297617892301
## 107626      2.09893260798678 -0.188503357422022  1.84799787870675
## 100861837  60.0354369270688   1.76258722180297  0.413367940819585
## 100861924  70.3085529839158 -0.757018298034826  0.470176666394451
## 170942      407.572590388916 -0.321253809935283  0.267531015042773
##           stat          pvalue          padj
##           <numeric>       <numeric>       <numeric>
## 497097      5.85081897183044 4.89158399055336e-09 1.52128680547954e-08
## 19888       0                      1                      1
## 20671      3.97868298186363 6.92980676409098e-05 0.000153603252478534
## 27395      -4.8998947864588 9.58879908891501e-07 2.52323667976689e-06
## 18777      -1.15295995413986 0.24892682179196  0.319245926805013
## ...
## ...
## 100862004 -0.592335046737235 0.553626251818812 0.62671677703091
## 107626     -0.102004098378045 0.918753422019503 0.940801019911589
## 100861837  4.26396691119366 2.00829197245236e-05 4.70326352425135e-05
## 100861924 -1.61007202641513 0.107382133154954 0.151258614218578
## 170942     -1.20080959541801 0.229825068881452 0.297801946926946

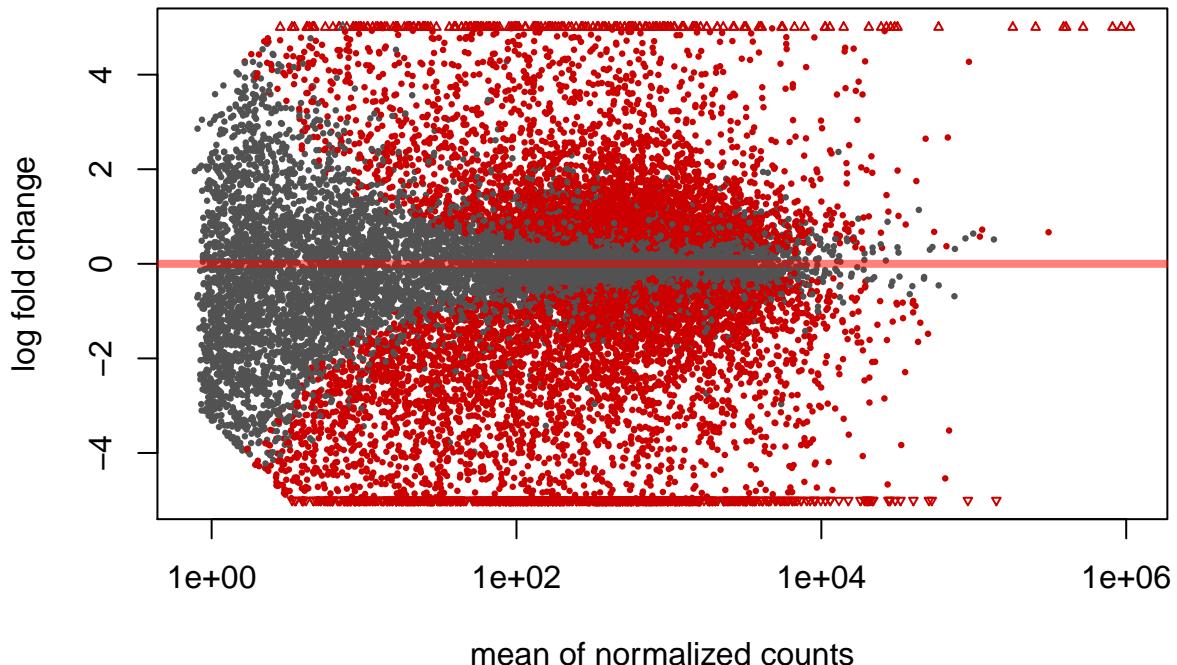
```

## Visualizing results

### MA-plot

In DESeq2, the function `plotMA` shows the log<sub>2</sub> fold changes attributable to a given variable over the mean of normalized counts for all the samples in the `DESeqDataSet`.

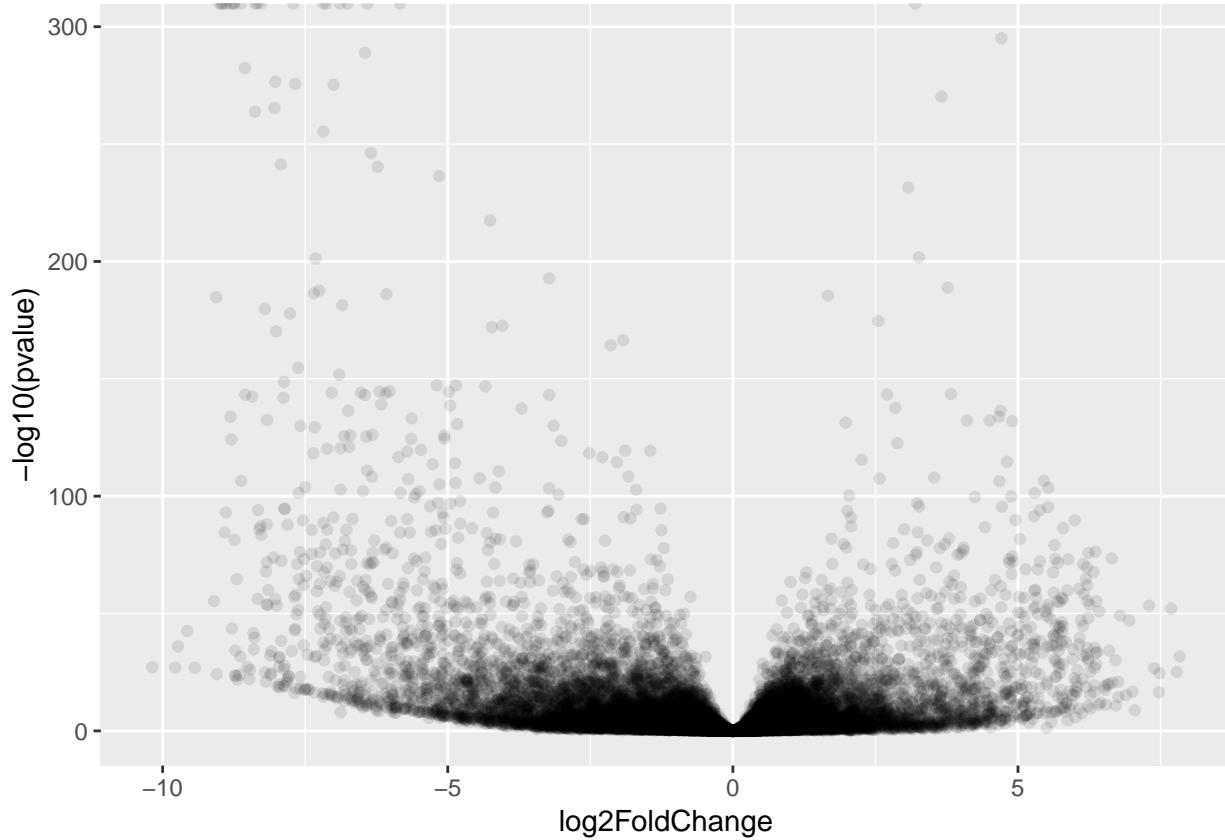
```
plotMA(res, ylim=c(-5, 5))
```



### Volcano plot

Volcano plots allow to visualize p-values and fold-change:

```
df<-as.data.frame(res@listData)
ggplot(df, aes(x=log2FoldChange, y=-log10(pvalue))) + geom_point(alpha = 0.1)
```



## GSEA

Download the gene set file from here. Several sets can be dowloaded. We choose the Hallmark gene sets with entrez gene ids and save it in `./data`. You should now have a file `h.all.v6.2.symbols.gmt` in the `./data` folder. Let's read it:

```
pathwaysH <- read.csv("./data/h.all.v6.2.symbols.gmt", sep="\t", header=F, stringsAsFactor=FALSE)
```

We need to do a little formatting before using the `fgsea()` function

```
#read the geneset file
# the format is different than when we do: load("human_H_v5p2.rdata")
# this causes a problem.
genesets <- lapply( 1:nrow(pathwaysH), function(x){
  row <- pathwaysH[x, ]
  current_string = paste( unlist( row[ 3:length( row)], use.names = FALSE), collapse="\t")
  current_string <- gsub("[\t]+$","", current_string)
  current_string <- strsplit(current_string, split="\t")[[1]]
  return(current_string)
})
names(genesets) <- pathwaysH[[1]]
```

We also need to rank genes. We choose to rank them according to the log2 fold-change stored in our `res` object.

```

ranks <- data.frame(ID = rownames(res), log2FC = res$logFoldChange)
idx_match <- match(rownames(res), annot$ENTREZID)
ranks$symbol <- toupper(annot$SYMBOL[idx_match])
ranks <- ranks[order(ranks$log2FC, decreasing = TRUE), ]

```

We keep only distinct gene symbols (the one with highest FC is kept)

```

ranks <- distinct(ranks, symbol, .keep_all = TRUE)
stats <- ranks$log2FC
names(stats) <- ranks$symbol

```

Run GSEA

```

library(fgsea)
fgseaRes <- fgsea(pathways = genesets, stats = stats, minSize=15, maxSize = 500, nperm=1000)

```

Lets look at the top 10 most enriched gene sets.

```

topUp <- fgseaRes %>%
  filter(ES > 0) %>%
  arrange(pval) %>%
  head(10)

topUp %>% dplyr::select(pathway, pval, ES)

##                               pathway      pval      ES
## 1    HALLMARK_ESTROGEN_RESPONSE_EARLY 0.01612903 0.3274123
## 2    HALLMARK_G2M_CHECKPOINT 0.01666667 0.3604773
## 3    HALLMARK_E2F_TARGETS 0.01724138 0.4130443
## 4    HALLMARK_MYC_TARGETS_V1 0.01724138 0.3116729
## 5    HALLMARK_OXIDATIVE_PHOSPHORYLATION 0.01724138 0.4514836
## 6    HALLMARKADIPOGENESIS 0.03389831 0.2883684
## 7    HALLMARK_FATTY_ACID_METABOLISM 0.04464286 0.2974395
## 8    HALLMARK_ESTROGEN_RESPONSE_LATE 0.04761905 0.2809479
## 9    HALLMARK_PROTEIN_SECRETION 0.04964539 0.3219071
## 10   HALLMARK_HEME_METABOLISM 0.10810811 0.2714692

```

and the top 10 most depleted gene sets.

```

topDown <- fgseaRes %>%
  filter(ES < 0) %>%
  arrange(pval) %>%
  head(10)

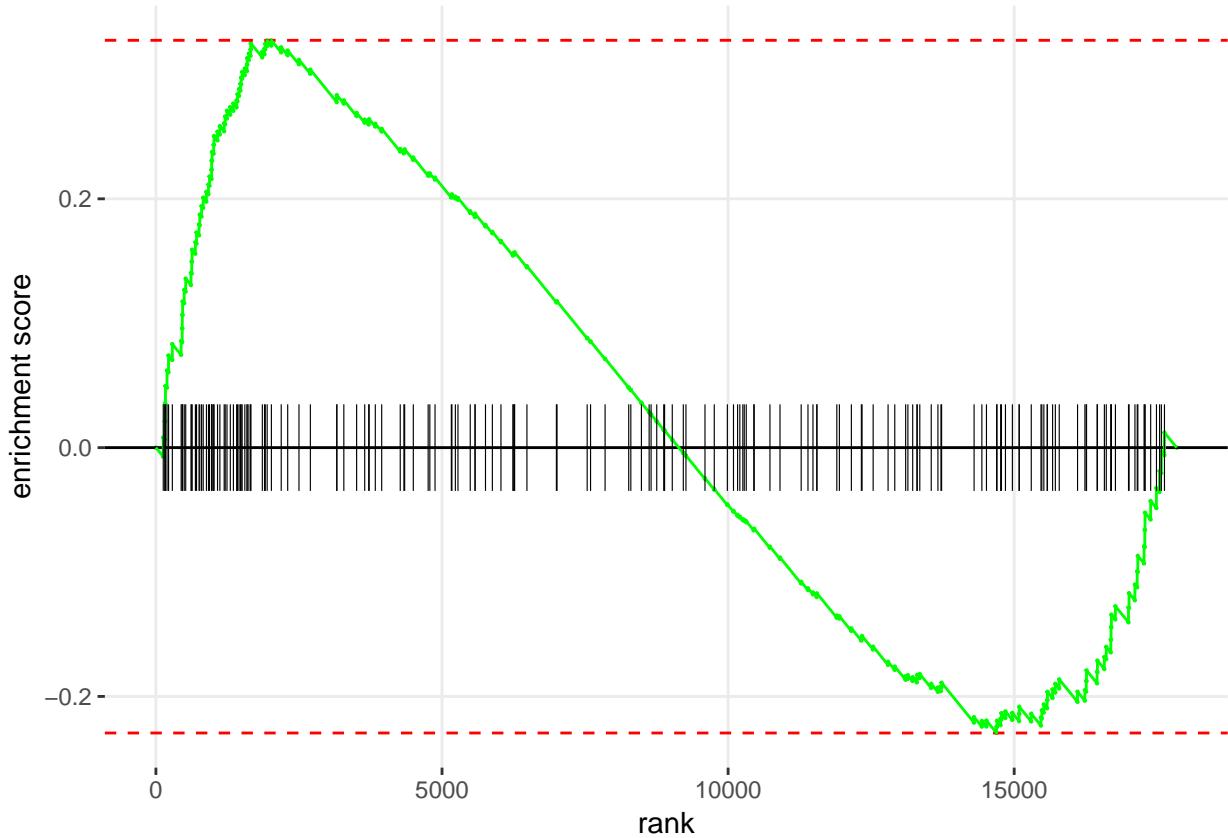
topDown %>% dplyr::select(pathway, pval, ES)

##                               pathway      pval      ES
## 1  HALLMARK_EPITHELIAL_MESENCHYMAL_TRANSITION 0.001059322 -0.6936600
## 2  HALLMARK_APICAL_JUNCTION 0.001064963 -0.5079610
## 3  HALLMARK_MYOGENESIS 0.001066098 -0.5528161
## 4  HALLMARK_UV_RESPONSE_DN 0.001117318 -0.5144255
## 5  HALLMARK_HYPOXIA 0.002127660 -0.4325774
## 6  HALLMARK_WNT_BETA_CATENIN_SIGNALING 0.004010695 -0.5816659
## 7  HALLMARK_COAGULATION 0.011337868 -0.4363982
## 8  HALLMARK_HEDGEHOG_SIGNALING 0.013812155 -0.5907059
## 9  HALLMARK_ANGIOGENESIS 0.024623803 -0.5525925
## 10 HALLMARK_KRAS_SIGNALING_DN 0.025698324 -0.4058336

```

Let's see how the genes within a pathway are distributed along the ranked genes.

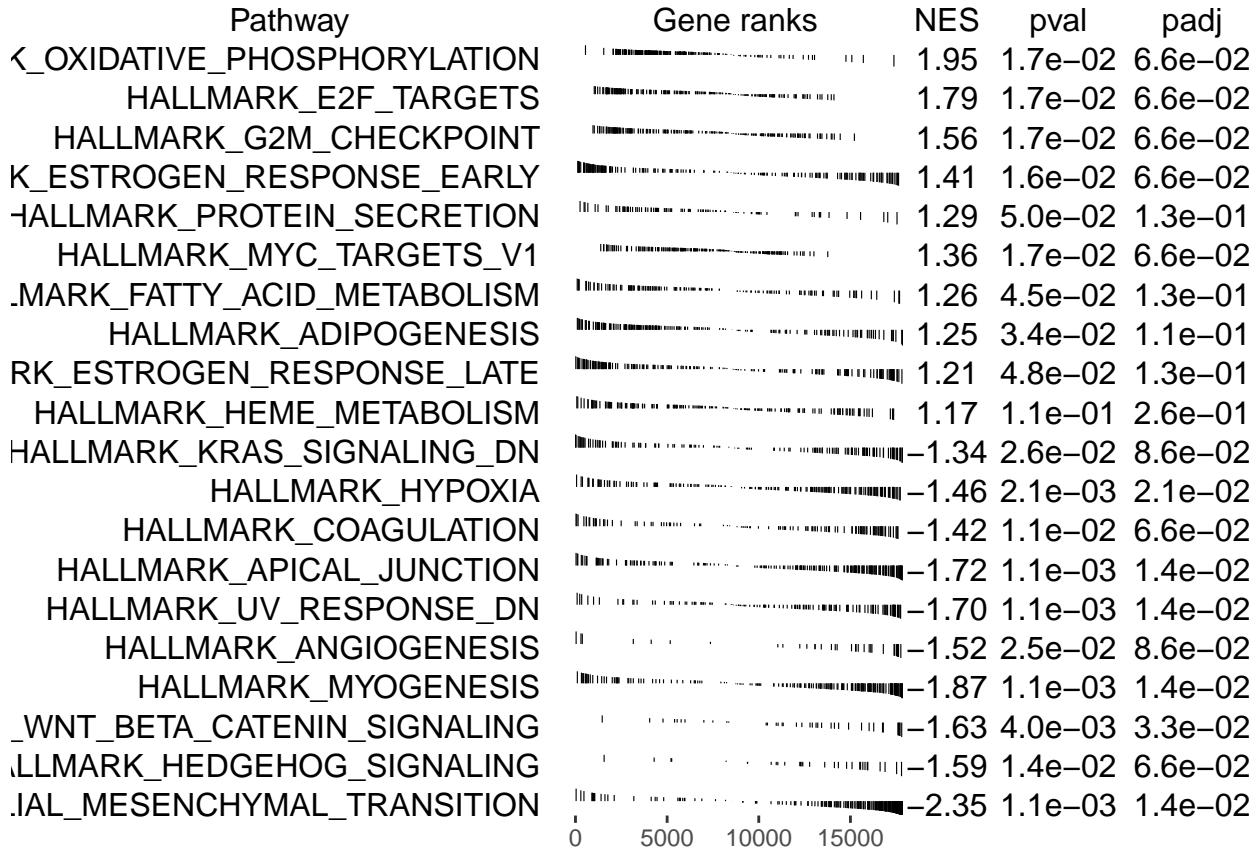
```
plotEnrichment(genesets[[topUp$pathway[1]]], stats)
```



We can also look at several pathways at once

```
topPathways <- bind_rows(topUp, topDown) %>%
  arrange(-ES)

plotGseaTable(genesets[topPathways$pathway],
              stats,
              fgseaRes,
              gseaParam = 0.5)
```



Re-run the analysis on another gene set of your choice.

## Additional Material

### Normalization

Normalization is performed using `estimateSizeFactors`. This function first normalize rows by the geometric mean and returns the median value for each column.

```
dds.norm <- estimateSizeFactors(dds)
sizeFactors(dds)

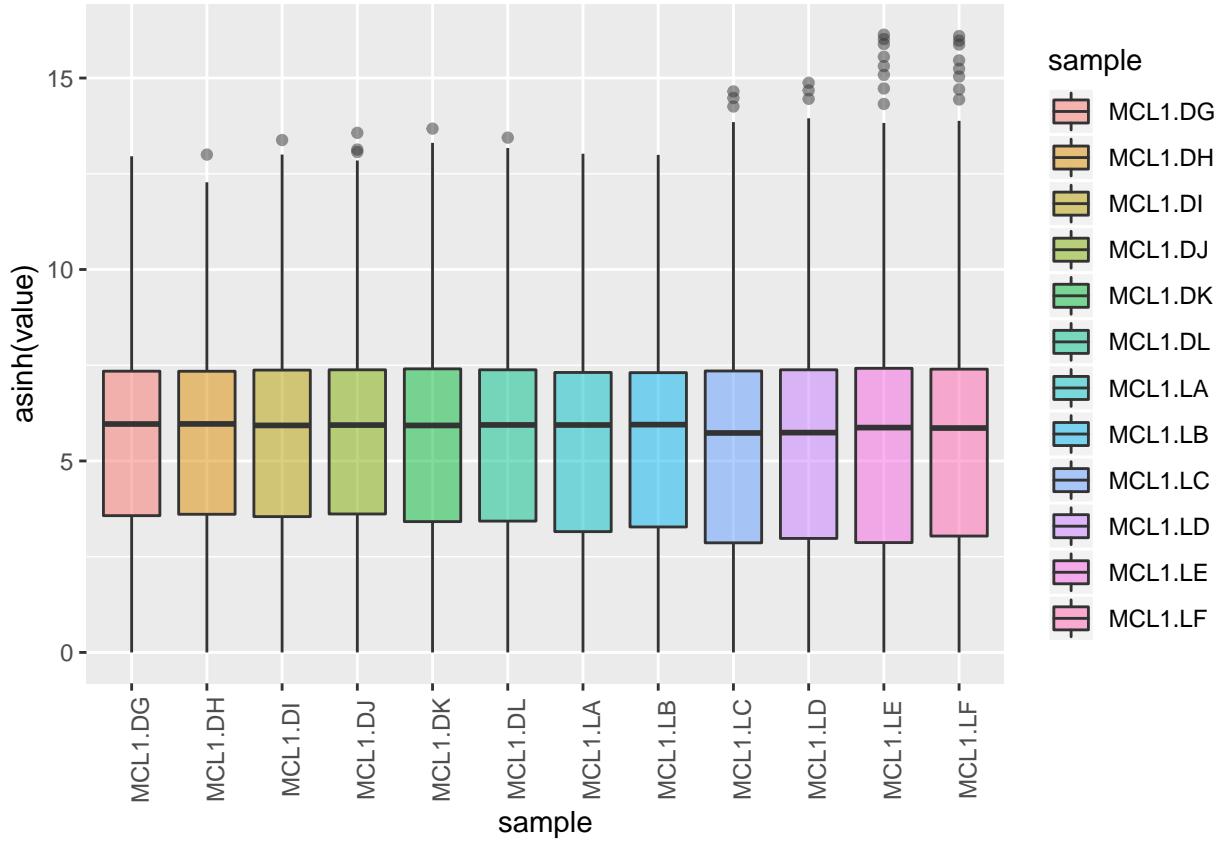
##   MCL1.DG   MCL1.DH   MCL1.DI   MCL1.DJ   MCL1.DK   MCL1.DL   MCL1.LA
## 1.2934903 1.1933709 1.2113142 1.0768623 0.9859083 0.9729893 1.2822154
##   MCL1.LB   MCL1.LC   MCL1.LD   MCL1.LE   MCL1.LF
## 1.3615404 1.0315620 0.9190180 0.5697346 0.5758505
```

Normalized data can now be retrieved using `counts(dds.norm, normalized=TRUE)`. Let's see what the normalized data looks like:

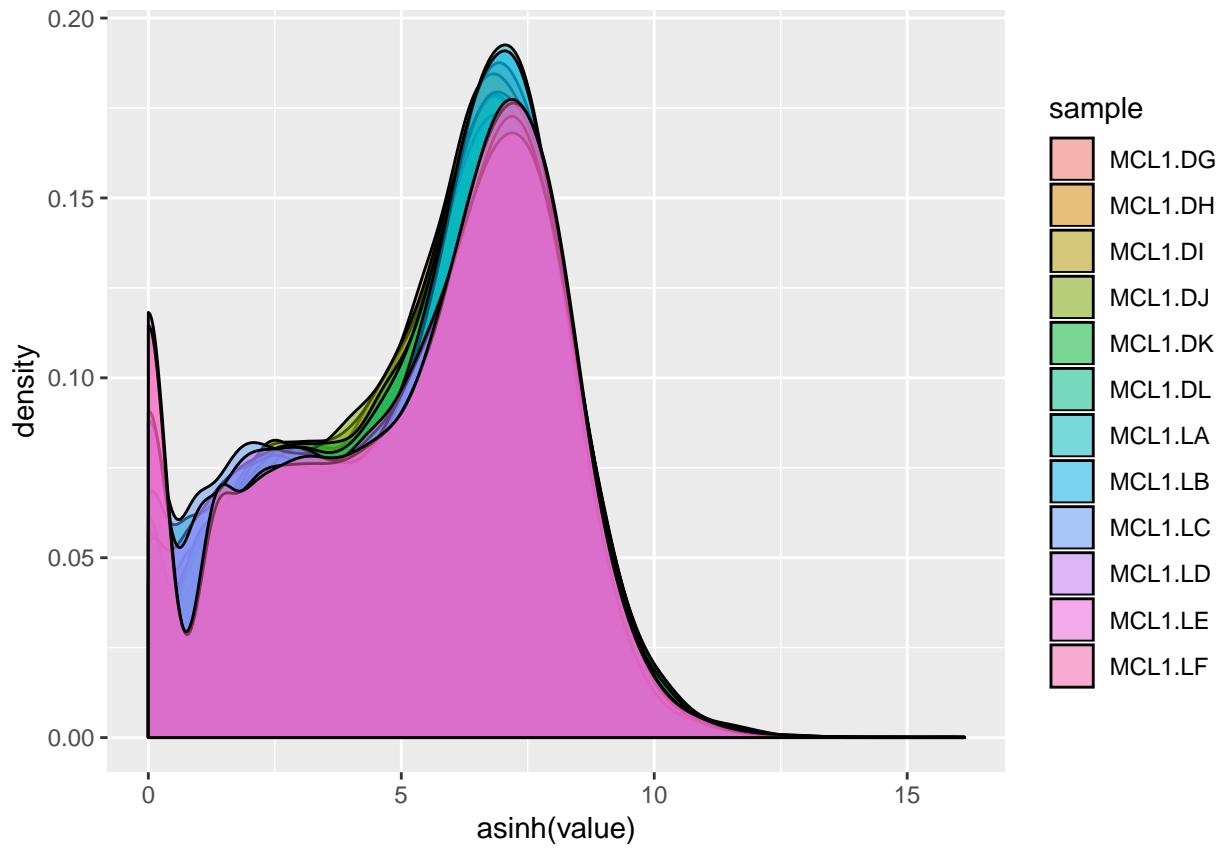
```
library(reshape2)
library(ggplot2)
df_norm <- as.data.frame(counts(dds.norm, normalized=TRUE))
df_norm [[ "GeneID" ]] <- rownames(df_norm )
df_melt_norm <- melt(df_norm , id.vars = "GeneID")
df_melt_norm <- dplyr::rename(df_melt_norm, sample = variable)

ggplot(df_melt_norm, aes(x=sample, y = asinh(value), fill = sample)) +
  theme(axis.text.x = element_text(angle = 90)) +
```

```
geom_boxplot(alpha = 0.5)
```



```
ggplot(df_melt_norm, aes(x = asinh(value), fill = sample)) +  
  geom_density(alpha = 0.5)
```



### Mean-variance relationship

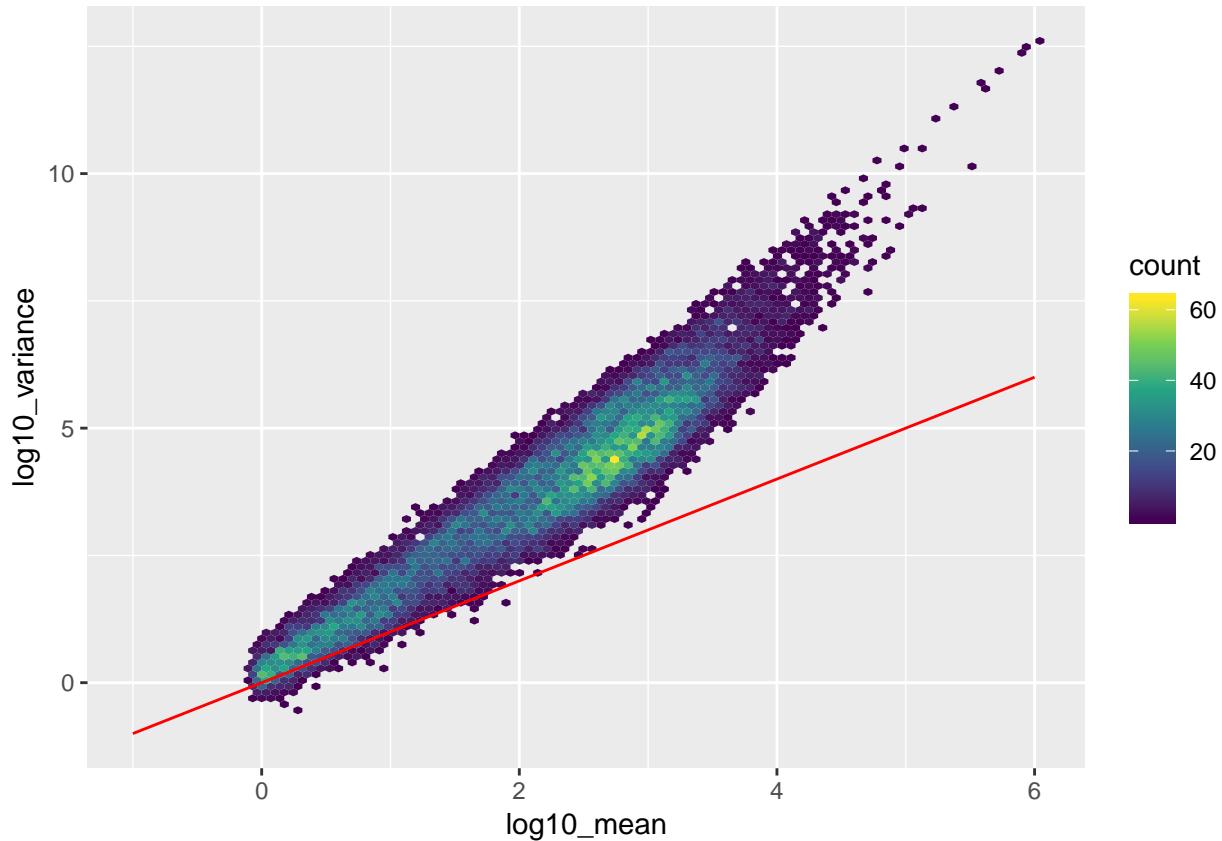
As you can see from the following plot the relationship between variance and mean is not linear as would be the case for a Poisson distribution (red line).

```
## Computing mean and variance
df_norm <- as.data.frame( counts(dds.norm, normalized=TRUE) )
mean <- apply(df_norm, 1, mean, na.rm = TRUE)
variance <- apply(df_norm, 1, var, na.rm = TRUE)

df_norm$log10_mean <- log10(mean)
df_norm$log10_variance <- log10(variance)

p1<-ggplot(df_norm, aes(x=log10_mean, y=log10_variance)) +
  geom_hex(bins=100) +
  scale_fill_viridis_c()+
  annotate("segment", x = -1, y = -1, xend = 6, yend=6, color = "red")
```

p1



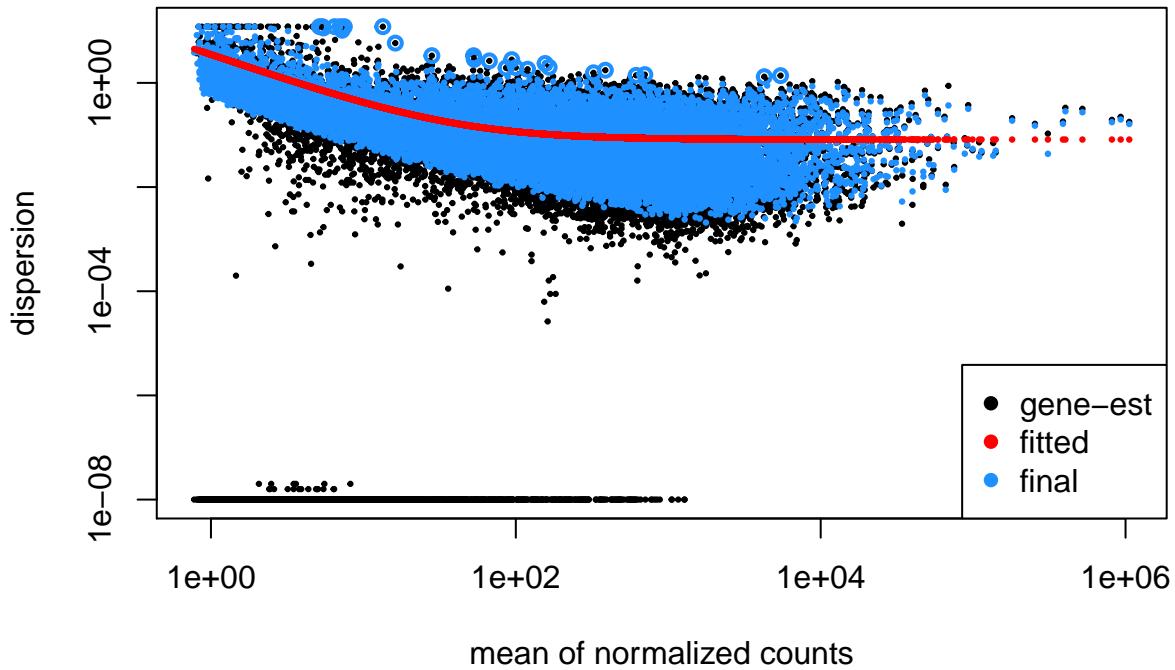
### Modelling read counts through a negative binomial

To perform differential expression call DESeq will assume that, for each gene, the read counts are generated by a negative binomial distribution. One problem here will be to estimate, for each gene, the two parameters of the negative binomial distribution: mean and dispersion.

-The mean will be estimated from the observed normalized counts in both conditions. -The first step will be to compute a gene-wise dispersion. When the number of available samples is insufficient to obtain a reliable estimator of the variance for each gene, DESeq will apply a shrinkage strategy, which assumes that counts produced by genes with similar expression level (counts) have similar variance (note that this is a strong assumption). DESeq will regress the gene-wise dispersion onto the means of the normalized counts to obtain an estimate of the dispersion that will be subsequently used to build the binomial model for each gene. Performing estimation of dispersion parameter and display a diagnostic plot

```
## Performing estimation of dispersion parameter
dds.disp <- estimateDispersions(dds.norm)

## A diagnostic plot which
## shows the mean of normalized counts (x axis)
## and dispersion estimate for each genes
plotDispEts(dds.disp)
```



### Performing differential expression

Now that a negative binomial model has been fitted for each gene, the `nbinomWaldTest` can be used to test for differential expression.

```
dds.wald <- nbinomWaldTest(dds.disp)
resultsNames(dds.wald)
```

```
## [1] "Intercept"           "CellType_luminal_vs_basal"
## [3] "Status_pregnant_vs_lactate" "Status_virgin_vs_lactate"
```

As before, we can see the result table. We keep only results with adjusted p-values (correction for multiple tests computed with the Benjamini–Hochberg procedure) below 0.01 using `alpha`.

```
res.wald <- results(dds.wald, alpha=0.01, pAdjustMethod="BH", name = "CellType_luminal_vs_basal")
summary(res.wald)
```

```
##
## out of 18178 with nonzero total read count
## adjusted p-value < 0.01
## LFC > 0 (up)      : 4946, 27%
## LFC < 0 (down)    : 5579, 31%
## outliers [1]       : 0, 0%
## low counts [2]     : 0, 0%
## (mean count < 1)
## [1] see 'cooksCutoff' argument of ?results
## [2] see 'independentFiltering' argument of ?results
head(res.wald)
```

```
## log2 fold change (MLE): CellType luminal vs basal
## Wald test p-value: CellType luminal vs basal
## DataFrame with 6 rows and 6 columns
##          baseMean    log2FoldChange        lfcSE
##          <numeric>    <numeric>    <numeric>
```

```

## 497097 126.481585433975 -10.1820437902145 0.930694407261515
## 19888 1.95697709143844 3.067152056161 1.23908077258863
## 20671 50.5979096011351 -2.71203870035955 0.337521822867245
## 27395 358.719720532762 0.799404687904274 0.127971895175
## 18777 744.356910502922 0.106722300367076 0.134155732874136
## 21399 1097.91010559132 -0.361177585114161 0.0647349104833116
##           stat              pvalue          padj
##             <numeric>          <numeric>      <numeric>
## 497097 -10.9402653661305 7.39830425247185e-28 1.00288124311285e-26
## 19888  2.47534472652113  0.0133107627365738  0.0215288766816833
## 20671 -8.03515066765403 9.34639809275225e-16 6.19842482780191e-15
## 27395  6.24672071013012 4.19159576750842e-10 1.74080027100224e-09
## 18777  0.795510546442337 0.42631657857844  0.488097690615066
## 21399 -5.57933242538849 2.41443422519718e-08 8.38387494663502e-08

```

## References

Fu, Nai Yang, Anne C Rios, Bhupinder Pal, Rina Soetanto, Aaron T L Lun, Kevin Liu, Tamara Beck, et al. 2015. “EGF-mediated induction of Mcl-1 at the switch to lactation is essential for alveolar cell survival.” *Nature Cell Biology* 17 (4): 365–75. doi:10.1038/ncb3117.